

A Probabilistic Model Revealing Shortcomings in Lua's Hybrid Tables

Pablo Rotondo

LIGM, Université Gustave Eiffel

Joint work with

Conrado Martínez (UPC), Cyril Nicaud (LIGM)

COCOON 2022,
Online, 23 October, 2022.

Introduction

- ▶ Aim: study and model actual implementations
 - Engineers sometimes choose innovative implementations
e.g., TimSort in Python.
 - Study choices in depth, make recommendations.

Introduction

- ▶ Aim: study and model actual implementations
 - Engineers sometimes choose innovative implementations e.g., TimSort in Python.
 - Study choices in depth, make recommendations.
- ▶ The Lua programming language
 - Scripting language widely used in the gaming industry,
 - Efficient, lightweight (few Kb of C code!), embeddable.

Introduction

- ▶ Aim: study and model actual implementations

- Engineers sometimes choose innovative implementations
e.g., TimSort in Python.
- Study choices in depth, make recommendations.

- ▶ The Lua programming language

- Scripting language widely used in the gaming industry,
- Efficient, lightweight (few Kb of C code!), embeddable.

⇒ Lua 5.0 introduced several innovations,
among them a new **Table** structure.

Introduction: table structure in Lua

- ▶ Only data-structuring mechanism in Lua
 - assignment $H[x]=y$, any types of x and y .

Introduction: table structure in Lua

- ▶ Only data-structuring mechanism in Lua
 - assignment $H[x]=y$, any types of x and y .
- ▶ Implementation
 - originally a simple hash-table up to Lua 4.

Introduction: table structure in Lua

- ▶ Only data-structuring mechanism in Lua
 - assignment $H[x]=y$, any types of x and y .
- ▶ Implementation
 - originally a simple hash-table up to Lua 4.
 - Lua 5 introduced a hybrid table-array,

Introduction: table structure in Lua

- ▶ Only data-structuring mechanism in Lua
 - assignment $H[x]=y$, any types of x and y .
- ▶ Implementation
 - originally a simple hash-table up to Lua 4.
 - Lua 5 introduced a hybrid table-array,
 - integer keys form array $\{1, \dots, n\}$ at least half-full.

Introduction: table structure in Lua

- ▶ Only data-structuring mechanism in Lua
 - assignment $H[x]=y$, any types of x and y .
- ▶ Implementation
 - originally a simple hash-table up to Lua 4.
 - Lua 5 introduced a hybrid table-array,
 - integer keys form array $\{1, \dots, n\}$ at least half-full.

In our work we

- ▶ study the hash-table mechanism, [*main result*]

Introduction: table structure in Lua

- ▶ Only data-structuring mechanism in Lua
 - assignment $H[x]=y$, any types of x and y .
- ▶ Implementation
 - originally a simple hash-table up to Lua 4.
 - Lua 5 introduced a hybrid table-array,
 - integer keys form array $\{1, \dots, n\}$ at least half-full.

In our work we

- ▶ study the hash-table mechanism, [*main result*]
- ▶ in worst case, but most importantly,
we introduce a reasonable probabilistic model.

Introduction: table structure in Lua

- ▶ Only data-structuring mechanism in Lua
 - assignment $H[x]=y$, any types of x and y .
- ▶ Implementation
 - originally a simple hash-table up to Lua 4.
 - Lua 5 introduced a hybrid table-array,
 - integer keys form array $\{1, \dots, n\}$ at least half-full.

In our work we

- ▶ study the hybrid table mechanism, [*main result*]
- ▶ in worst case, but most importantly,
we introduce a reasonable probabilistic model.
- ▶ present also an analysis of the hybrid table-array.

Introduction: worst case Lua hash-table

Lua's hash-table aims for space-efficiency:

Introduction: worst case Lua hash-table

Lua's hash-table aims for space-efficiency:

- ▶ insertions and lookups work in amortized $O(1)$
even if table is full.

Introduction: worst case Lua hash-table

Lua's hash-table aims for space-efficiency:

- ▶ insertions and lookups work in amortized $O(1)$
even if table is full.
- ▶ but we show there is a degradation if deletions are allowed.

Introduction: worst case Lua hash-table

Lua's hash-table aims for space-efficiency:

- ▶ insertions and lookups work in amortized $O(1)$
even if table is full.
- ▶ but we show there is a degradation if deletions are allowed.

Consider sequences of T insertions/deletions starting from an empty table

Introduction: worst case Lua hash-table

Lua's hash-table aims for space-efficiency:

- ▶ insertions and lookups work in amortized $O(1)$ even if table is full.
- ▶ but we show there is a degradation if deletions are allowed.

Consider sequences of T insertions/deletions starting from an empty table

Proposition: worst case

There is sequence of operations giving time $\Theta(T^2)$.

Introduction: worst case Lua hash-table

Lua's hash-table aims for space-efficiency:

- ▶ insertions and lookups work in amortized $O(1)$ even if table is full.
- ▶ but we show there is a degradation if deletions are allowed.

Consider sequences of T insertions/deletions starting from an empty table

Proposition: worst case

There is sequence of operations giving time $\Theta(T^2)$.

- ▶ The example requires an unlikely cycle of delete-insert.

Introduction: worst case Lua hash-table

Lua's hash-table aims for space-efficiency:

- ▶ insertions and lookups work in amortized $O(1)$
even if table is full.
- ▶ but we show there is a degradation if deletions are allowed.

Consider sequences of T insertions/deletions starting from an empty table

Proposition: worst case

There is sequence of operations giving time $\Theta(T^2)$.

- ▶ The example requires an unlikely cycle of delete-insert.
- ▶ Trouble for more realistic examples ?

Introduction: probabilistic model for the hash-table

Simple Probabilistic model

Consider $p > \frac{1}{2}$. A random sequence of T insertion/deletions:



- ▶ with probability p insert a new element,
- ▶ with probability $1 - p$ delete an element.

Introduction: probabilistic model for the hash-table

Simple Probabilistic model

Consider $p > \frac{1}{2}$. A random sequence of T insertion/deletions:

- ▶ with probability p insert a new element,
- ▶ with probability $1 - p$ delete an element.



Main result: Lua hash-table



With high probability time is $\Omega(T \log T)$.

Introduction: probabilistic model for the hash-table

Simple Probabilistic model

Consider $p > \frac{1}{2}$. A random sequence of T insertion/deletions:

- ▶ with probability p insert a new element,
- ▶ with probability $1 - p$ delete an element.



Main result: Lua hash-table

With high probability time is $\Omega(T \log T)$.



Plan of the talk

1. The Lua hashmap
2. The probabilistic model
3. Conclusions and further work

The *pure* hashmap mechanism

Lua's *hashmap* consists of

- ▶ an array H of size $M = 2^m$,

The *pure* hashmap mechanism

Lua's *hashmap* consists of

- ▶ an array H of size $M = 2^m$,
- ▶ a hash function $h(x) = x \bmod M$, [*we do not discuss choice of h*]

The *pure* hashmap mechanism


Lua's hashmap consists of

- ▶ an array H of size $M = 2^m$,
- ▶ a hash function $h(x) = x \bmod M$, [*we do not discuss choice of h*]
- ▶ entries: *key*, *value*, *index of next entry* in chain

The *pure* hashmap mechanism

Lua's *hashmap* consists of

- ▶ an array H of size $M = 2^m$,
- ▶ a hash function $h(x) = x \bmod M$, [we do not discuss choice of h]
- ▶ entries: *key*, *value*, *index of next entry* in chain



The *pure* hashmap mechanism: insertion

Insertions work as follows: key x

- ▶ if position $h(x)$ free \Rightarrow insert

The *pure* hashmap mechanism: insertion

Insertions work as follows: key x

- ▶ if position $h(x)$ free \Rightarrow insert
- ▶ else position $h(x)$ is occupied by key y ,
 - if $h(y) = h(x) \Rightarrow$ put x into a **free position**, update chain $\text{pos}(y) \rightarrow \text{pos}(z) \rightarrow \dots$ to $\text{pos}(y) \rightarrow \text{pos}(x) \rightarrow \text{pos}(z) \rightarrow \dots$

The *pure* hashmap mechanism: insertion

Insertions work as follows: key x

- ▶ if position $h(x)$ free \Rightarrow insert
- ▶ else position $h(x)$ is occupied by key y ,
 - if $h(y) = h(x) \Rightarrow$ put x into a **free position**, update chain $\text{pos}(y) \rightarrow \text{pos}(z) \rightarrow \dots$ to $\text{pos}(y) \rightarrow \text{pos}(x) \rightarrow \text{pos}(z) \rightarrow \dots$
 - if $h(y) \neq h(x) \Rightarrow$ we migrate y into a **free position**, updating its chain and put x at position $h(x)$.


	0	1	2	3	4	5	6	7
H					a 11 nil			

$$h(a) = 4, h(b) = 4,$$

The *pure* hashmap mechanism: insertion

Insertions work as follows: key x

- ▶ if position $h(x)$ free \Rightarrow insert
- ▶ else position $h(x)$ is occupied by key y ,
 - if $h(y) = h(x) \Rightarrow$ put x into a **free position**, update chain $\text{pos}(y) \rightarrow \text{pos}(z) \rightarrow \dots$ to $\text{pos}(y) \rightarrow \text{pos}(x) \rightarrow \text{pos}(z) \rightarrow \dots$
 - if $h(y) \neq h(x) \Rightarrow$ we migrate y into a **free position**, updating its chain and put x at position $h(x)$.




$$h(a) = 4, h(b) = 4, h(d) = 4,$$

The *pure* hashmap mechanism: insertion

Insertions work as follows: key x

- ▶ if position $h(x)$ free \Rightarrow insert
- ▶ else position $h(x)$ is occupied by key y ,
 - if $h(y) = h(x) \Rightarrow$ put x into a **free position**, update chain $\text{pos}(y) \rightarrow \text{pos}(z) \rightarrow \dots$ to $\text{pos}(y) \rightarrow \text{pos}(x) \rightarrow \text{pos}(z) \rightarrow \dots$
 - if $h(y) \neq h(x) \Rightarrow$ we migrate y into a **free position**, updating its chain and put x at position $h(x)$.




$$h(a) = 4, h(b) = 4, h(d) = 4, h(e) = 7,$$

The *pure* hashmap mechanism: insertion

Insertions work as follows: key x

- ▶ if position $h(x)$ free \Rightarrow insert
- ▶ else position $h(x)$ is occupied by key y ,
 - if $h(y) = h(x) \Rightarrow$ put x into a **free position**, update chain $\text{pos}(y) \rightarrow \text{pos}(z) \rightarrow \dots$ to $\text{pos}(y) \rightarrow \text{pos}(x) \rightarrow \text{pos}(z) \rightarrow \dots$
 - if $h(y) \neq h(x) \Rightarrow$ we migrate y into a **free position**, updating its chain and put x at position $h(x)$.




$$h(a) = 4, h(b) = 4, h(d) = 4, h(e) = 7, h(f) = 7$$

The *pure* hashmap mechanism: insertion

Insertions work as follows: key x

- ▶ if position $h(x)$ free \Rightarrow insert
- ▶ else position $h(x)$ is occupied by key y ,
 - if $h(y) = h(x) \Rightarrow$ put x into a **free position**, update chain $\text{pos}(y) \rightarrow \text{pos}(z) \rightarrow \dots$ to $\text{pos}(y) \rightarrow \text{pos}(x) \rightarrow \text{pos}(z) \rightarrow \dots$
 - if $h(y) \neq h(x) \Rightarrow$ we migrate y into a **free position**, updating its chain and put x at position $h(x)$.




$$h(a) = 4, h(b) = 4, h(d) = 4, h(e) = 7, h(f) = 7$$

The *pure* hashmap mechanism: insertion

Insertions work as follows: key x

- ▶ if position $h(x)$ free \Rightarrow insert
- ▶ else position $h(x)$ is occupied by key y ,
 - if $h(y) = h(x) \Rightarrow$ put x into a **free position**, update chain $\text{pos}(y) \rightarrow \text{pos}(z) \rightarrow \dots$ to $\text{pos}(y) \rightarrow \text{pos}(x) \rightarrow \text{pos}(z) \rightarrow \dots$
 - if $h(y) \neq h(x) \Rightarrow$ we migrate y into a **free position**, updating its chain and put x at position $h(x)$.




$$h(a) = 4, h(b) = 4, h(d) = 4, h(e) = 7, h(f) = 7$$

The *pure* hashmap mechanism: insertion

Insertions work as follows: key x

- ▶ if position $h(x)$ free \Rightarrow insert
- ▶ else position $h(x)$ is occupied by key y ,
 - if $h(y) = h(x) \Rightarrow$ put x into a **free position**, update chain $\text{pos}(y) \rightarrow \text{pos}(z) \rightarrow \dots$ to $\text{pos}(y) \rightarrow \text{pos}(x) \rightarrow \text{pos}(z) \rightarrow \dots$
 - if $h(y) \neq h(x) \Rightarrow$ we migrate y into a **free position**, updating its chain and put x at position $h(x)$.



Finding a free position: use pointer starting at end and moving to the left. If pointer exits, we **rehash**.

The *pure* hashmap mechanism: deletion

Deletions are simple:

- ▶ the *value* is marked as *nil*,

The *pure* hashmap mechanism: deletion

Deletions are simple:

- ▶ the *value* is marked as *nil*,
- ▶ chaining (*next cell*) kept intact.

The *pure* hashmap mechanism: deletion

Deletions are simple:

- ▶ the *value* is marked as **nil**,
- ▶ chaining (*next cell*) kept intact.

Deleted spot y

- ▶ can be reused to insert x when $h(x) = y$,

The *pure* hashmap mechanism: deletion

Deletions are simple:

- ▶ the *value* is marked as **nil**,
- ▶ chaining (*next cell*) kept intact.

Deleted spot y

- ▶ can be reused to insert x when $h(x) = y$,
- ▶ not taken into account by **free position pointer**
 \implies necessary to keep previous chaining

The *pure* hashmap mechanism: deletion

Deletions are simple:

- ▶ the *value* is marked as **nil**,
- ▶ chaining (*next cell*) kept intact.

Deleted spot y

- ▶ can be reused to insert x when $h(x) = y$,
- ▶ not taken into account by **free position pointer**
 \implies necessary to keep previous chaining

... **deleted spots** are cleaned up during **rehashing**

The *pure* hashmap mechanism: rehashing

Finding a free position: use pointer starting at end and moving to the left. If pointer exits, we **rehash**.

- ▶ Hashtable is then full, maybe with deleted cells.

The *pure* hashmap mechanism: rehashing

Finding a free position: use pointer starting at end and moving to the left. If pointer exits, we **rehash**.

- ▶ Hashtable is then full, maybe with deleted cells.
- ▶ Count actual used cells n ,

The *pure* hashmap mechanism: rehashing

Finding a free position: use pointer starting at end and moving to the left. If pointer exits, we **rehash**.

- ▶ Hashtable is then full, maybe with deleted cells.
- ▶ Count actual used cells n ,
- ▶ New hashtable of size $M = 2^m$, smallest m s.t. $n + 1 \leq 2^m$,
 $\implies +1$ for inserted element.

The *pure* hashmap mechanism: rehashing

Finding a free position: use pointer starting at end and moving to the left. If pointer exits, we **rehash**.

- ▶ Hashtable is then full, maybe with deleted cells.
- ▶ Count actual used cells n ,
- ▶ New hashtable of size $M = 2^m$, smallest m s.t. $n + 1 \leq 2^m$,
 $\implies +1$ for inserted element.

Worst-case scenario

- ▶ insert until filling hashtable of size $M = 2^m$,
- ▶ alternate M deletion/insertions,
- ▶ insertions induce rehash unless deleted cell is picked,
- ▶ complexity $\Theta(M^2)$ for $3M$ operations.

The *pure* hashmap mechanism: rehashing

Finding a free position: use pointer starting at end and moving to the left. If pointer exits, we **rehash**.

- ▶ Hashtable is then full, maybe with deleted cells.
- ▶ Count actual used cells n ,
- ▶ New hashtable of size $M = 2^m$, smallest m s.t. $n + 1 \leq 2^m$,
 $\implies +1$ for inserted element.

Worst-case scenario

- ▶ insert until filling hashtable of size $M = 2^m$,
- ▶ alternate M deletion/insertions,
- ▶ insertions induce rehash unless deleted cell is picked,
- ▶ complexity $\Theta(M^2)$ for $3M$ operations.

... but it is not very realistic

The probabilistic model

We set a more interesting yet simple model

Probabilistic model

Fix $p > \frac{1}{2}$ and apply T insertion/deletions from an empty table:

- ▶ with probability p insert a new element,
- ▶ with probability $1 - p$ delete an element **among present ones**.

The probabilistic model

We set a more interesting yet simple model

Probabilistic model

Fix $p > \frac{1}{2}$ and apply T insertion/deletions from an empty table:

- ▶ with probability p insert a new element,
- ▶ with probability $1 - p$ delete an element **among present ones**.

Hashtable tends to **grow**: $\# \text{ keys} \approx pT - (1 - p)T = (2p - 1)T$

The probabilistic model

We set a more interesting yet simple model

Probabilistic model

Fix $p > \frac{1}{2}$ and apply T insertion/deletions from an empty table:

- ▶ with probability p insert a new element,
- ▶ with probability $1 - p$ delete an element **among present ones**.

Hashtable tends to **grow**: $\# \text{ keys} \approx pT - (1 - p)T = (2p - 1)T$

Theorem (Martínez, Nicaud, R 2022)

With high probability, Lua uses $\Omega(T \log T)$ time for this process.

- ▶ **Intuition:** Large number of rehashes
that serve only to remove few **nil** cells.

The probabilistic model

We set a more interesting yet simple model

Probabilistic model

Fix $p > \frac{1}{2}$ and apply T insertion/deletions from an empty table:

- ▶ with probability p insert a new element,
- ▶ with probability $1 - p$ delete an element **among present ones**.

Hashtable tends to **grow**: $\# \text{ keys} \approx pT - (1 - p)T = (2p - 1)T$

Theorem (Martínez, Nicaud, R 2022)

With high probability, Lua uses $\Omega(T \log T)$ time for this process.

- ▶ **Intuition**: Large number of **reshashes**
that serve only to remove few **nil** cells.
- ▶ Each **rehash** costs linear time $\Theta(M)$.



The probabilistic model: proof sketch

Theorem (Martínez, Nicaud, R 2022)

With high probability, Lua uses $\Omega(T \log T)$ time for this process.

- ▶ Number of keys in **hashmap** after t operations $\approx (2p - 1)t$,
with high probability size M **only increases**

The probabilistic model: proof sketch

Theorem (Martínez, Nicaud, R 2022)

With high probability, Lua uses $\Omega(T \log T)$ time for this process.

- ▶ Number of keys in **hashmap** after t operations $\approx (2p - 1)t$,
with high probability size M **only increases**
- ▶ First time we rehash into size $M = 2^m$ of order $\Theta(T)$:
 $\implies \# \text{keys} = 2^{m-1} + 1$ and $2^{m-1} - 1$ free spots.

The probabilistic model: proof sketch

Theorem (Martínez, Nicaud, R 2022)

With high probability, Lua uses $\Omega(T \log T)$ time for this process.

- ▶ Number of keys in **hashmap** after t operations $\approx (2p - 1)t$,
with high probability size M **only increases**
- ▶ First time we rehash into size $M = 2^m$ of order $\Theta(T)$:
 $\implies \# \text{keys} = 2^{m-1} + 1$ and $2^{m-1} - 1$ free spots.
- ▶ But then random deletions slow down growth of M :
 \implies from f free spots, we obtain γf after next rehash.

The probabilistic model: proof sketch

Theorem (Martínez, Nicaud, R 2022)

With high probability, Lua uses $\Omega(T \log T)$ time for this process.

- ▶ Number of keys in **hashmap** after t operations $\approx (2p - 1)t$,
with high probability size M **only increases**
- ▶ First time we rehash into size $M = 2^m$ of order $\Theta(T)$:
 $\implies \# \text{keys} = 2^{m-1} + 1$ and $2^{m-1} - 1$ free spots.
- ▶ But then random deletions slow down growth of M :
 \implies from f free spots, we obtain γf after next rehash.

Lemma

If the hashmap has size M and just after a rehash it contains $f \gg \sqrt{M}$ free spots, then at the next rehash it still has size M and contains at least γf free spots (whp).

The probabilistic model: proof sketch

Theorem (Martínez, Nicaud, R 2022)

With high probability, Lua uses $\Omega(T \log T)$ time for this process.

- ▶ Number of keys in **hashmap** after t operations $\approx (2p - 1)t$,
with high probability size M **only increases**
- ▶ First time we rehash into size $M = 2^m$ of order $\Theta(T)$:
 $\implies \# \text{keys} = 2^{m-1} + 1$ and $2^{m-1} - 1$ free spots.
- ▶ But then random deletions slow down growth of M :
 \implies from f free spots, we obtain γf after next rehash.

Lemma

If the hashmap has size M and just after a rehash it contains $f \gg \sqrt{M}$ free spots, then at the next rehash it still has size M and contains at least γf free spots (whp).

... at least $\log M$ rehashes to increase M

The probabilistic model: proof sketch

With (very) high probability:

- ▶ the hashtable is never empty after $t = 0$,
- ▶ we rehash at some point.

The probabilistic model: proof sketch

With (very) high probability:

- ▶ the hashtable is never empty after $t = 0$,
- ▶ we rehash at some point.

Under these conditions, between two rehashes, the **number of deleted cells** satisfies the recurrence (starting from $\delta_{t_0} = 0$)

$$\delta_{t+1} = \begin{cases} \delta_t - 1 & \text{with probability } \frac{p\delta_t}{M} & [\textit{insertion at deleted key}], \\ \delta_t & \text{with probability } p\left(1 - \frac{\delta_t}{M}\right) & [\textit{insertion at free cell}], \\ \delta_t + 1 & \text{with probability } 1 - p & [\textit{deletion}]. \end{cases}$$

The probabilistic model: proof sketch

$$\delta_{t+1} = \begin{cases} \delta_t - 1 & \text{with probability } \frac{p\delta_t}{M} & [\textit{insertion at deleted key}], \\ \delta_t & \text{with probability } p\left(1 - \frac{\delta_t}{M}\right) & [\textit{insertion at free cell}], \\ \delta_t + 1 & \text{with probability } 1 - p & [\textit{deletion}]. \end{cases}$$

- **Equilibrium point** at $\delta_t \approx \frac{1-p}{p}M$,
 - ⊗ when $\delta_t < \frac{1-p}{p}M$ tendency to increase,
 - ⊗ when $\delta_t > \frac{1-p}{p}M$ tendency to decrease,

The probabilistic model: proof sketch

$$\delta_{t+1} = \begin{cases} \delta_t - 1 & \text{with probability } \frac{p\delta_t}{M} & [\textit{insertion at deleted key}], \\ \delta_t & \text{with probability } p\left(1 - \frac{\delta_t}{M}\right) & [\textit{insertion at free cell}], \\ \delta_t + 1 & \text{with probability } 1 - p & [\textit{deletion}]. \end{cases}$$

- ▶ **Equilibrium point** at $\delta_t \approx \frac{1-p}{p}M$,
 - ⊗ when $\delta_t < \frac{1-p}{p}M$ tendency to increase,
 - ⊗ when $\delta_t > \frac{1-p}{p}M$ tendency to decrease,
- ▶ **Rehash** occurs a fraction before reaching equilibrium
... sufficiently before to make δ_t increase linearly

The probabilistic model: proof sketch

$$\delta_{t+1} = \begin{cases} \delta_t - 1 & \text{with probability } \frac{p\delta_t}{M} & [\textit{insertion at deleted key}], \\ \delta_t & \text{with probability } p\left(1 - \frac{\delta_t}{M}\right) & [\textit{insertion at free cell}], \\ \delta_t + 1 & \text{with probability } 1 - p & [\textit{deletion}]. \end{cases}$$

- ▶ **Equilibrium point** at $\delta_t \approx \frac{1-p}{p}M$,
 - ⊗ when $\delta_t < \frac{1-p}{p}M$ tendency to increase,
 - ⊗ when $\delta_t > \frac{1-p}{p}M$ tendency to decrease,
- ▶ **Rehash** occurs a fraction before reaching equilibrium
... sufficiently before to make δ_t increase linearly
- ▶ **Tools:** concentration inequalities.

Hybrid Tables and insertions

- ⊗ Rehashes into same size hashtables pile up to $\Omega(T \log T)$.

Hybrid Tables and insertions

- ⊗ Rehashes into same size hashtables pile up to $\Omega(T \log T)$.

And without deletions?

- ⊗ Problem arises when considering effects of deletions.

Hybrid Tables and insertions

- ⊗ Rehashes into same size hashtables pile up to $\Omega(T \log T)$.

And without deletions?

- ⊗ Problem arises when considering **effects of deletions**.
- ⊗ The **hybrid data-structure** presents a similar issue:
using the array-part “simulates” deletions on the hash-part

Proposition [only insertions]

Inserting n elements into Lua's table takes $\Theta(n \log n)$ in the worst case.

Hybrid Tables and insertions

- ⊗ Reshapes into same size hashtables pile up to $\Omega(T \log T)$.

And without deletions?

- ⊗ Problem arises when considering **effects of deletions**.
- ⊗ The **hybrid data-structure** presents a similar issue:
using the array-part “simulates” deletions on the hash-part

Proposition [only insertions]

Inserting n elements into Lua's table takes $\Theta(n \log n)$ in the worst case.

Example: inserting $-(2^k - 1), -(2^k - 2), \dots, -1, 0, 1, \dots, 2^k$

Recap and conclusions

- ⊗ Lua's **hybrid data-structure** is an interesting idea.
- ⊗ We have presented a **simple and natural probabilistic model**
revealing shortcomings in Lua's hashtables.

Recap and conclusions

- ⊗ Lua's **hybrid data-structure** is an interesting idea.
- ⊗ We have presented a **simple and natural probabilistic model**
revealing shortcomings in Lua's hashtables.
- ⊗ Issue can be fixed by ensuring **more room when rehashing**.
- ⊗ This would also fix the hybrid part.

Recap and conclusions

- ⊗ Lua's **hybrid data-structure** is an interesting idea.
- ⊗ We have presented a **simple and natural probabilistic model** revealing shortcomings in Lua's hashtables.
- ⊗ Issue can be fixed by ensuring **more room when rehashing**.
- ⊗ This would also fix the hybrid part.

Conclusions

- ⊗ Will Lua conceptors take this into account?
- ⊗ Important to model and study algorithms implemented in practice.

Thank you!